



Vulnerability Detection - Sanitizers

Holistic Software Security

Aravind Machiry



Bug Manifestations

- How bugs affect program behavior?
 - If we have exhaustive test cases:
 - Actual output != Expected output.
 - In the absence of test cases, i.e., Fuzzing:
 - Memory errors: Program Crashes (SIGSEGV) => Access/Execute invalid memory.
 - There could be bugs which do not result in SIGSEGV.



Silent Bugs

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```

Here, if $i == 2$ (off by one error), the program may not crash?

- Why?

Runtime Stack

Return Address
i
j
unsigned a[2]

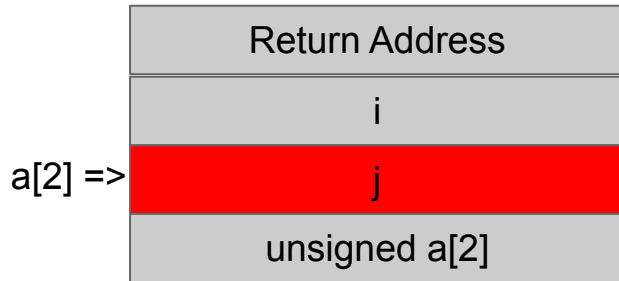
Silent Bugs

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```

Here, if $i == 2$ (off by one error), the program may not crash?

- Why?

Runtime Stack





Improving bug detection

- The behavior of a bug, especially memory corruption, depends on the program state and execution environment.
- Can we detect these bugs without relying on program state?
 - Fuzzing: we detect a bug if it results in the program crash (SIGSEGV).
 - Idea: Make all bugs result in program crashes.



Sanitizers

- Change the program such that *we detect bugs when they occur* instead of waiting for the bugs to result in crash.
- Mechanism: Instrument the program by adding additional checks for detecting bugs.

Sanitizers: Overview

Original Program



Instrumentation

Instrumented Program



Program with
additional Checks

Original Program

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```



Array out-of-bounds Sanitizer

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    if (i < 2) {  
        a[i] = j;  
    } else { CRASH}  
    ...  
    return 0;  
}
```

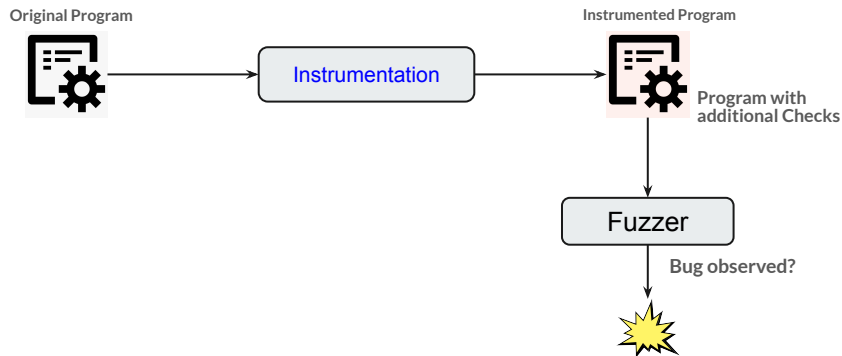
Instrumented Program



Real world Sanitizers

- Usually bug specific. Examples:
 - MemorySanitizer: Detects *uninitialized reads*.
 - AddressSanitizer: Detects *invalid memory accesses*.
- **General instrumentation idea:** At all instructions in the program where the bug can occur, add a check to detect the bug.
 - AddressSanitizer: Detects Invalid Memory Accesses.
 - Invalid Memory access can occur at load and store instructions.
 - Instrument every load and store to check if the used address is invalid (i.e., does not belong to a program object).

Sanitizers: Usage



Original Program

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```

Array out-of-bounds Sanitizer

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    if (i < 2) {  
        a[i] = j;  
    } else { CRASH}  
    ...  
    return 0;  
}
```

Instrumented Program



Why can't we always use sanitizers?

They detect bugs at runtime => Why can't we just use sanitizers and not worry about bugs, as they will never lead to vulnerabilities.

**THERE'S
NO
FREE
LUNCH**



Sanitizers introduce **a lot** of overhead.



Sanitizers Implementation

- Sanitizers need to maintain lot of additional state to check for the possibility of bugs.
 - AddressSanitizer: Detects Invalid Memory Accesses:
 - Need to **maintain metadata** regarding which memory (i.e., address) is valid v/s invalid.
 - Tricky: Handling dynamic memory allocation.

Popular research direction: Smart and efficient way to maintain metadata.

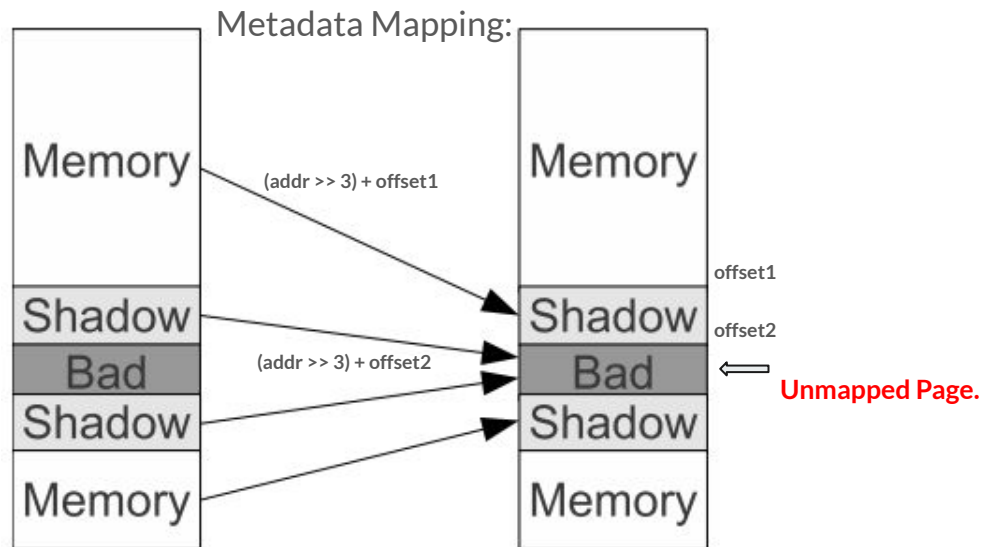


AddressSanitizer (ASan)

- Metadata (or shadow memory):
 - One eighth of the virtual memory will be used to maintain metadata:
 - One bit of metadata for each byte of application memory.
 - Bit is zero: The corresponding address is valid else invalid.
- Accessing metadata for a given address (Addr):
 - Direct Mapping:

```
// Checking 8-byte access
MetadataAddr = (Addr >> 3) + Offset;
if (*MetadataAddr != 0)
    ReportAndCrash(Addr);
```

ASan: Mapping





ASan: Usage

a.c

```
void foo(T *a) {  
    *a = 0x1234;  
}
```



clang -fsanitize=address a.c -c -DT=long



```
push %rax  
mov %rdi,%rax  
shr $0x3,%rax  
mov $0x1000000000000,%rcx  
or %rax,%rcx  
cmpb $0x0,(%rcx) # Compare Shadow with 0  
jne 23 <foo+0x23> # To Error  
movq $0x1234,(%rdi) # Original store  
pop %rax  
retq  
callq __asan_report_store8 # Error
```



ASan: Conclusion

- One of the most popular sanitizers: Used extensively in fuzzing.
- Overhead:
 - Adds additional instructions:
 - Memory overhead: **~3X** (Consumes thrice the amount of memory).
 - Slowdown: **~2X** (Runs at half the speed).



ThreadSanitizer

- Detects data races.
- Where can data races happen i.e., which instructions it should track?
- How to detect a data race? What metadata should be maintained?



Other sanitizers (supported by clang)

- -fsanitize=address: AddressSanitizer, a memory error detector.
- -fsanitize=thread: ThreadSanitizer, a data race detector.
- -fsanitize=memory: MemorySanitizer, a detector of uninitialized reads. Requires instrumentation of all program code.
- -fsanitize=undefined: UndefinedBehaviorSanitizer, a fast and compatible undefined behavior checker.
- -fsanitize=dataflow: DataFlowSanitizer, a general data flow analysis.



Sanitizers: Final Thoughts

- They increase the ability of fuzzing to find bugs.
- Always use them with fuzzers: Performance impact does not matter much - lets throw more machines.
- New sanitizers => Always appreciated and could have a high impact.
- Decreasing overhead of sanitizers: Appreciated but may have less impact.