

# Assignment 0

Ricardo Andres Calvo Mendez

## Problem 1

### Exercise A

```
int main(int argc, char **argv) {  
    char buf[10];  
    strcpy(buf, argv[0]);  
    ....  
}
```

#### Analysis

Here, the program copies the command-line argument ( `argv[0]` is always the program name) into the buffer `buf` without checking its size.

#### Failing scenario:

When `strlen(argv[0]) > 10`, the buffer overflow occurs.

### Exercise B

```
size_t s;  
char *p;  
scanf("%lu", &s);  
p = (char*)malloc(s + 4);  
if (p) {  
    strcpy(p, "HDR");  
    fgets(p + 3, s, stdin);  
} else {  
    printf("Out of memory!\n");  
    return -1;  
}
```

#### Analysis

Here, the function `malloc()` is used to allocate memory of size `s + 4`, where `s` is provided by user input. The program does not validate if `s + 4` overflows the maximum value representable by `size_t`. If an overflow occurs, `malloc()`

would allocate less memory than intended, leading to a buffer overflow when `fgets()` is called.

### Failing scenario:

When the user provides as input a value between  $2^{64} - 4$  and  $2^{64} - 1$  (assuming `size_t` is 64 bits)

## Excercise C

```
static void webize(char *str, char *dfstr, int dfsize) {
    char *cp1;
    char *cp2;
    for (cp1 = str, cp2 = dfstr; *cp1 != '\0' && cp2 - dfstr
        < dfsize - 1; ++cp1, ++cp2) {
        switch (*cp1) {
            case '<':
                *cp2++ = '&';
                *cp2++ = 'l';
                *cp2++ = 't';
                *cp2 = ';';
                break;
            case '>':
                *cp2++ = '&';
                *cp2++ = 'g';
                *cp2++ = 't';
                *cp2 = ';';
                break;
            default:
                *cp2 = *cp1;
                break;
        }
        *cp2 = '\0';
    }
}
```

### Analysis

Here, the function `webize()` copies characters from the input string `str` to the output buffer `dfstr`, replacing `<` and `>` with their HTML entity equivalents. However, the function does not properly account for the increased length of the output string when these replacements occur. Each `<` or `>` character is replaced by a 4-character sequence (`&lt;` or `&gt;`), which can lead to a buffer overflow if there are enough such characters in the input string.

### Failing scenario:

Let  $a :=$  number of `'<'` and `'>'` characters in `str`.

When `strlen(str) + 3a ≥ dfsize`, the buffer overflow occurs.

## Excercise D

```
int *p;
int q[20];
unsigned s;
...
memset(q, 0, sizeof(q));
...
p = malloc(s);
if (p != NULL){
    memset(p, 'A', sizeof(p));
} else {
    return -1;
}
```

### Analysis

Here, The problem is in the second `memset()` call. the return value of `sizeof(p)` is not the size of the allocated memory, it is the size of the pointer itself (constant 8 bytes).

### Failing scenario:

When `s ≤ 8`, the buffer overflow occurs.

## Exercise E

```
char format[20];
// Read format to display the log string.
scanf("%19s", format);
...
// Print the log_str in required format.
printf(format, log_str);
...
```

### Analysis

Here, The problem is the `format` input string, the user can introduce a format string to print out a given amount of bytes, probably exceeding the allocated buffer size for `log_str`.

### Failing scenario:

When `format` is intended to print more bytes than `log_str` can hold.

## Exercise F

```
class base {
    public:
        base() {
        }
        ~base() {
        }
}

class sub : public base {
    public:
        sub() {
        }
        ~sub() {
        }
}

int main() {
    base *b = new sub();
    ....
    delete b;
}
```

### Analysis

Here, There is a parent class `base` and a child class `sub`. The issue arises when a pointer of type `base` is used to delete an object of type `sub`, which can lead to resource leaks because the destructor of the child class is not called.

To fix the problem either `b` should be of type `sub` or the base class destructor should be made virtual.

## Exercise G

```
char fl;
....
int ret = sscanf(buf, %s, &fl);
if (ret != 1) {
    printf("Read Error\n");
    return -1;
}
```

### Analysis

Here, the problem is that the size of `fl` is not being checked before using it in `sscanf`. This can lead to buffer overflows if the input exceeds the allocated

size.

**Failing scenario:**

When `fl` points to a buffer smaller than the input being read.

## Problem 2

A. If we avoid storing return address on runtime stack then stack-based buffer overflows do not cause any security issues (especially, control-flow hijacking)

**Answer**

No, In general that no prevent ANY security issues, It prevents only control-flow hijacking attacks because that addresses cannot be used to redirect execution flow.

B. We can always prove that a given program does not have any security vulnerabilities.

**Answer**

As demonstrated in class, the answer to this question is the same as the halting problem: it is undecidable in general.

C. Exhaustive testing proves that the a given program does not have any bugs.

**Answer**

No, In general that falls in the halting problem again, you can prove most of the inputs but not all of them.

## Problem 3

**A. A process can know physical addresses of its virtual addresses. Justify your answer in either case.**

**Answer**

No, the virtual address is converted to physical address by Operating system. The program does not have knowledge of physical addresses. Only kernel-like processes can access physical addresses directly.

**B. A process can read and write memory that belong to the operating system kernel. Justify your answer in either case.**

### Answer

No, the operating system prevents processes from accessing memory that does not belong to them. The typical error that occurs when a process tries to access memory outside its allocated space is a segmentation fault, which results in the termination of the process.

**C. Operating system should always sanitize (i.e., verify) addresses given by a user process. Why? E.g., Destination address provided for read/write syscall.**

### Answer

Yes, the operating system should always ensure that user programs operate within their allocated memory space and do not interfere with each other's memory.

**D. Is there any security issue in the following code? Justify your answer in either case.**

```
unsigned gl;
char flag_buf[4];
...
unsigned i;
if (!copy_from_user(&i, buf, sizeof(i))) {
    if (i < 4) {
        if (!copy_from_user(&gl, buf, sizeof(gl))) {
            flag_buf[gl] = 0;
        }
    }
}
...
```

### Answer

Yes, a buffer overflow can occur when it is accessing using `gl` as index `flag_buf[gl] = 0`. `gl` is a value read from the user (I assume `copy_from_user()` is used to read this value from a human). the value of `gl` is not being validated