

# Preliminaries

Ricardo Andres Calvo Mendez

PUID: 36688820 · rcalvome@purdue.edu

ECE 69500 · Hardware Software Security

September 3, 2026

Santiago Torres Arias

---

## 1 Installing the Required Tools

I use Arch Linux as my operating system. I use `pacman` to manage packages.

I found that the package names listed in the assignment correspond directly to packages in `pacman`. I used the `--needed` flag to avoid reinstalling packages if they were already present on my system.

### Terminal

```
→ ricardo ~ sudo pacman -S --needed arm-none-eabi-gcc
→ ricardo ~ sudo pacman -S --needed arm-none-eabi-binutils
→ ricardo ~ sudo pacman -S --needed qemu-system-arm
```

Now, I ran the `-ql` flag on each of them to see which files those packages brought to my computer. I realized that the output was super long, so I redirected it to a file.

### Terminal

```
→ ricardo ~ pacman -Ql arm-none-eabi-gcc > arm-none-eabi-gcc-files.txt
→ ricardo ~ pacman -Ql arm-none-eabi-binutils > arm-none-eabi-binutils-files.txt
→ ricardo ~ pacman -Ql qemu-system-arm > qemu-system-arm-files.txt
```

**For `arm-none-eabi-gcc`:** It contains 2,969 entries. Using the folders, I grouped the contents like this:

1. The path `/usr/arm-none-eabi/include/c++/` contains  $\sim 1,650$  entries ( $\sim 56\%$ ) consisting mainly of C++ headers.
2. The path `/usr/arm-none-eabi/lib/` contains  $\sim 250$  entries ( $\sim 9\%$ ) related to the required libraries for the package.
3. The path `/usr/bin/` contains  $\sim 20$  entries ( $\sim 1\%$ ) related to the executables. I assume this because it is included in my `$PATH`.
4. The path `/usr/lib/gcc/arm-none-eabi/` contains  $\sim 1,000$  entries (35%) related to GCC drivers.
5. The last path `/usr/share/man/` contains manual pages for the installed commands.

For **arm-none-eabi-binutils**: It contains 84 entries.

1. The path `/usr/arm-none-eabi/bin/` contains  $\sim 10$  entries ( $\sim 12\%$ ) consisting mainly of binaries.
2. The path `/usr/arm-none-eabi/lib/` contains  $\sim 20$  entries ( $\sim 24\%$ ) related to the required libraries for the package.
3. The path `/usr/bin/` contains  $\sim 15$  entries ( $\sim 18\%$ ) related to the executables.
4. The path `/usr/share/man/` contains the remaining files ( $\sim 46\%$ ) for manual pages.

For **qemu-system-arm**: It contains only 16 entries.

1. The executable in the path `/usr/bin/`.
2. Some licenses in the path `/usr/share/licenses/`.
3. Manual pages at the path `/usr/share/man/`.

## 2 qemu-system-arm

Now, for this part of the assignment, I tried to run the command provided to run the “hello world,” but I got this message.

### Terminal

```
→ ricardo homework1 qemu-arm cpu-emu-test
zsh: command not found: qemu-arm
```

Checking with my trusty LLM, I realized that I was missing the package **qemu-user**, so I installed it using **pacman** as I did before.

### Terminal

```
→ ricardo homework1 sudo pacman -S --needed qemu-user
```

For **qemu-user**: It contains only 45 entries.

1. Some required tool executables in the path `/usr/bin/`.
2. Manual pages at the path `/usr/share/man/` as usual.

Now, it runs correctly!

### Terminal

```
→ ricardo homework1 qemu-arm cpu-emu-test
Hello world!
```

Continuing with the assignment, I ran the command to emulate a machine:

**Terminal**

```
→ ricardo ~ qemu-system-arm -machine mcimx6ul-evk
VNC server running on ::1:5900
```

I was expecting a monitor interface as shown in the reading. I again asked my trusty LLM for options to solve this problem.

I have two options to address this:

1. To install `qemu-ui-gtk` that would provide my system with a GTK interface to the emulated machine.
2. To add the flags `-display none -monitor stdio` so it will directly open the monitor in my terminal emulator.

I chose option 2 because I prefer to avoid unnecessary packages. Now, my updated command works correctly:

**Terminal**

```
→ ricardo ~ qemu-system-arm -machine mcimx6ul-evk -display none -monitor stdio

QEMU 11.1.1 monitor - type 'help' for more information
(qemu)
```

Now the setup is completed!

### 3 Building ARM binaries

Let's save the file `babysteps.s` with the recommended formatting:

**babysteps.s**

```
1  .global _start
2
3  _start:
4  top:
5      eor r1, r1
6      b top
```

To build this, I ran the commands recommended in the readings:

**Terminal**

```
→ ricardo homework1 arm-none-eabi-as -g babysteps.s -o babysteps.o
→ ricardo homework1 arm-none-eabi-ld -g babysteps.o -o babysteps.elf
```

After investigating this myself, I was able to understand what those commands do:

- The first command `arm-none-eabi-as` builds an object file for the ARM architecture. It reads the file `babysteps.s` and then the flag `-o` indicates the output path `babysteps.o`.
- The second command `arm-none-eabi-ld` links the object file assigning memory addresses and resolving symbols. It reads the file `babysteps.o` and then the flag `-o` indicates the output path `babysteps.elf`. The flag `-g` is just a flag for debugging tools.

Now, decomposing the assembly source code, I would describe it as follows:

- `.global _start` declares the symbol `_start` as global.
- `_start:` indicates that the symbol `_start` starts in that position.
- `top:` indicates that the symbol `top` starts in that position.
- `eor r1, r1` indicates an XOR operation. The evaluation of  $r1 \oplus r1$  is assigned to `r1`.
- `b top` indicates a jump to the `top:` symbol.

I would summarize the assembly as the following “equivalent” `c` code that is more human-readable:

`babysteps.c`

```
1  int start(){
2      int r1;
3      while(1){
4          r1 = r1 ^ r1;
5      }
6  }
```

It contains an endless loop with an assignment of an XOR operation.

## 4 Studying the compiled binary

I ran the command `readelf` as shown in the reading.

Terminal

```
→ ricardo homework1 readelf -h babysteps.elf
```

ELF Header:

```

Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00
Class:                               ELF32
Data:                                   2's complement, little endian
Version:                             1 (current)
OS/ABI:                               UNIX - System V
ABI Version:                          0
Type:                                   EXEC (Executable file)
Machine:                              ARM
Version:                              0x1
```

```

Entry point address:      0x8000
Start of program headers: 52 (bytes into file)
Start of section headers: 4580 (bytes into file)
Flags:                   0x5000200, Version5 EABI, soft-float ABI
Size of this header:      52 (bytes)
Size of program headers:  32 (bytes)
Number of program headers: 1
Size of section headers:  40 (bytes)
Number of section headers: 8
Section header string table index: 7

```

After investigating documentation, the command `readelf` displays all the information related to ELF files. ELF files are usually executable files in Unix systems.

Summarizing the information there:

- **Magic:** Traditional file magic number. The bytes `7f 45 4c 46` correspond to the ELF magic number.
- **Class:** 32-bit ELF format.
- **Data:** Little-endian representation.
- **Version:** ELF format version.
- **OS/ABI:** Operating system and ABI for which the file was generated.
- **ABI Version:** Version of the specified ABI.
- **Type:** ELF file type. `EXEC` means that it is an executable file.
- **Machine:** Target processor architecture.
- **Version:** ELF object file version.
- **Entry point address:** Address where program execution begins.
- **Start of program headers:** Offset from the beginning of the file to the program header table.
- **Start of section headers:** Offset from the beginning of the file to the section header table.
- **Flags:** Architecture-specific flags.
- **Size of this header:** Size of the ELF header.
- **Size of program headers:** Size of each entry in the program header table.
- **Number of program headers:** Number of entries in the program header table.
- **Size of section headers:** Size of each entry in the section header table.
- **Number of section headers:** Number of entries in the section header table.

- **Section header string table index:** Specifies which section header contains the string table used to store section names.

I ran the command `readelf` with the flag `--section-headers` to get more information.

#### Terminal

```
→ ricardo homework1 readelf --section-headers babysteps.elf
There are 8 section headers, starting at offset 0x11e4:
```

Section Headers:

| [Nr] | Name            | Type           | Addr     | Off    | Size   | ES | Flg | Lk | Inf | Al |
|------|-----------------|----------------|----------|--------|--------|----|-----|----|-----|----|
| [ 0] |                 | NULL           | 00000000 | 000000 | 000000 | 00 |     | 0  | 0   | 0  |
| [ 1] | .text           | PROGBITS       | 00008000 | 001000 | 000008 | 00 | AX  | 0  | 0   | 4  |
| [ 2] | .persistent     | PROGBITS       | 00009008 | 001008 | 000000 | 00 | WA  | 0  | 0   | 4  |
| [ 3] | .noinit         | NOBITS         | 00009008 | 000000 | 000000 | 00 | WA  | 0  | 0   | 4  |
| [ 4] | .ARM.attributes | ARM_ATTRIBUTES | 00000000 | 001008 | 000012 | 00 |     | 0  | 0   | 1  |
| [ 5] | .symtab         | SYMTAB         | 00000000 | 00101c | 000120 | 10 |     | 6  | 8   | 4  |
| [ 6] | .strtab         | STRTAB         | 00000000 | 00113c | 000062 | 00 |     | 0  | 0   | 1  |
| [ 7] | .shstrtab       | STRTAB         | 00000000 | 00119e | 000045 | 00 |     | 0  | 0   | 1  |

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), I (info),  
 L (link order), O (extra OS processing required), G (group), T (TLS),  
 C (compressed), x (unknown), o (OS specific), E (exclude),  
 D (mbind), y (purecode), p (processor specific)

Again, I researched this information and summarized it as follows:

- **Index 0:** Empty by convention.
- **.text:** Executable instructions of the program. It is allocated in memory and executable.
- **.persistent:** Data intended to persist across certain executions or resets. It is writable and allocated.
- **.noinit:** Uninitialized data that requires memory space but does not need to be stored in the ELF file. It is writable and allocated.
- **.ARM.attributes:** ARM-specific metadata describing architecture and ABI attributes.
- **.symtab:** The symbol table contains information about functions and variables.
- **.strtab:** String table containing the names of functions and variables.
- **.shstrtab:** The string table with the names of the sections.

## 5 Disassembling the compiled binary

Let's disassemble the ELF file we got.

## Terminal

```
→ ricardo homework1 arm-none-eabi-objdump -D babysteps.elf
```

```
babysteps.elf:      file format elf32-littlearm
```

```
Disassembly of section .text:
```

```
00008000 <_start>:
```

```
8000:      e0211001      eor      r1, r1, r1
8004:      eaffffffd      b        8000 <_start>
```

```
Disassembly of section .ARM.attributes:
```

```
00000000 <.ARM.attributes>:
```

```
0:  00001141      andeq   r1, r0, r1, asr #2
4:  61656100      cmnvs   r5, r0, lsl #2
8:  01006962      tsteq   r0, r2, ror #18
c:  00000007      andeq   r0, r0, r7
10: Address 0x10 is out of bounds.
```

According to the output we got, I see the disassembled code in the `.text` section. The output includes memory addresses, but its logic is the same as the source code we had before.

Cleaning up the disassembly output, I would recover the data like this:

## babysteps\_recovered.s

```
1  _start:
2      eor      r1, r1, r1
3      b        _start
```

I noticed the following points and I tried to give the best explanation I could:

1. **.global keyword is gone:** Investigating a little bit about this, the way to determine whether a label is global is to inspect the symbol table. Using the command `readelf -s babysteps.elf`, I was able to see it as global.
2. **Label `top` is gone:** In my experience working in compilers, it should've been removed by some sort of compiler optimization. It is clear that this label was not necessary due to the existence of `_start` at the same memory address.
3. **eor instruction now has 3 parameters:** This is because the original `eor r1, r1` is just a shortcut of the 3-parameter version. The destination register is the same as the first operand of the XOR operation.

Essentially, the changes are for optimization, but the semantics of the sources are equivalent.

## 6 Running ARM binaries we just built

To run the binaries in an emulated CPU, I crafted the following command.

### Terminal

```
→ ricardo homework1 qemu-arm babysteps.elf
(No output; stopped with Ctrl+C.)
```

I did not get any output because of the endless loop with no prints.

Now, let's run the binaries in an emulated machine, I crafted the following command.

### Terminal

```
→ ricardo homework1 qemu-system-arm
  -machine mcimx6ul-evk
  -kernel babysteps.elf
  -display none
  -monitor stdio
QEMU 11.1.1 monitor - type 'help' for more information
(qemu) stop
(qemu) info registers

CPU#0
R00=00000000 R01=00000000 R02=00000000 R03=00000000
R04=00000000 R05=00000000 R06=00000000 R07=00000000
R08=00000000 R09=00000000 R10=00000000 R11=00000000
R12=00000000 R13=00000000 R14=00000000 R15=00008000
...
FPSCR: 00000000
(qemu) quit
```

I removed part of the output to keep only the important information; I put ... where I cut the info.

The program didn't print any info (because it is not supposed to do so), but I can see that the register **R15=00008000** is moving in the endless loop.

To run this, I had to:

- Install the required packages in my Arch Linux installation using pacman.
  - The required compiler extensions.
  - Utils for binaries.
  - QEMU as the full hardware emulator (`qemu-system-arm`).
- Build a sample assembly `babysteps.s` and then run it in a full hardware emulation of the board `mcimx6ul-evk`.

Now, I investigated (with a little help from my agent) which board in the catalog of emulable boards is the closest to the `USB Armory MKII`.



I found the board `mcimx6ul-evk` as the closest for the following reasons:

- Both use processors from the `NXP i.MX6UL` family.
- Both use an `ARM Cortex-A7 CPU` core.
- The emulated board has peripherals and a memory layout similar to those of an embedded `i.MX6UL` system.

I will continue using this `mcimx6ul-evk` board from now on.

**i Extra Credit** Modify the assembly code to print a string to the `serial0` port upon startup.

For this, I modified the file as follows:

`babysteps_serial.s`

```

1  .equ UART1_UTXD, 0x02020040
2
3  .global _start
4
5  _start:
6      ldr r0, =message
7      ldr r1, =UART1_UTXD
8
9  print_character:
10     ldrb r2, [r0], #1
11     cmp r2, #0
12     beq .
13     str r2, [r1]
14     b print_character
15
16  message:
17     .asciz "Hello from serial0!\r\n"
```

More details about this:

- `.equ UART1_UTXD, 0x02020040` defines the address of the UART1 transmit register, which QEMU connects to the `serial0` output.
- `.global _start` makes the entry symbol globally visible.
- The first two `ldr` instructions store the address of the message in `r0` and the address of the UART transmit register in `r1`.
- `ldrb r2, [r0], #1` reads one byte of the message into `r2` and increments `r0` to point to the next character.
- `cmp r2, #0` and `beq .` detect the null byte at the end of the string and stay at the current instruction forever.

- `str r2, [r1]` writes the character to the UART transmit register, and `b print_character` repeats the process for the next character.
- `.asciz` stores the message with a null byte at the end; `\r\n` are a carriage return and a new line.

I then built and executed it.

#### Terminal

```
→ ricardo homework1 arm-none-eabi-as babysteps_serial.s -o babysteps_serial.o
→ ricardo homework1 arm-none-eabi-ld babysteps_serial.o -o babysteps_serial.elf
→ ricardo homework1 qemu-system-arm
    -machine mcimx6ul-evk
    -kernel babysteps_serial.elf
    -display none
    -monitor none
    -serial stdio
Hello from serial0!
```

## 7 Debugging our environment

Let's install the debugger.

#### Terminal

```
→ ricardo ~ sudo pacman -S --needed arm-none-eabi-gdb
```

Now, I ran the sample program in one terminal and the debugger in another terminal.

I had to replace the flag `-singlestep` with `-one-insn-per-tb` because the former is not available in the latest version of QEMU (QEMU 11.1.1).

#### Terminal 1

```
→ ricardo homework1 qemu-arm -one-insn-per-tb -g 1234 babysteps.elf
(Waiting for connection...)
```

#### Terminal 2

```
→ ricardo homework1 arm-none-eabi-gdb
GNU gdb (GDB) 17.2
Copyright (C) 2025 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "--host=x86_64-pc-linux-gnu --target=arm-none-eabi".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
```

```

<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
  <http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
(gdb) file babysteps.elf
Reading symbols from babysteps.elf...
(gdb) target remote localhost:1234
Remote debugging using localhost:1234
_start () at babysteps.s:5
5          eor r1, r1
(gdb) b babysteps.s:4
Breakpoint 1 at 0x8000: file babysteps.s, line 5.
(gdb) c
Continuing.

Breakpoint 1, _start () at babysteps.s:5
5          eor r1, r1

```

What I did so far is the following step-by-step:

1. I ran `babysteps.elf`, opening the port `1234` for debugging.
2. In the second terminal, I opened the debugger `GDB`.
3. Then, I used `file babysteps.elf` to load the symbols of that file.
4. Then, I connected to the emulated program via port `1234` of the localhost.
5. I set a breakpoint at the beginning of the endless loop; then the program stopped there.
6. Finally, I continued the execution, but the program hit the breakpoint again after going up in the loop.

## 8 Modifying our basic code

Now, let's create the memory checker described in the reading.

I crafted the following file.

```

memory_check.s
1  .global _start
2
3  _start:
4      ldr r0, =memory_start
5      ldr r1, =memory_end
6
7  check_memory:

```

```

8     ldr r2, [r0], #4
9     cmp r2, #0
10    bne explode
11    cmp r0, r1
12    bne check_memory
13
14    all_zero:
15        b .
16
17    explode:
18        b .
19
20    .data
21    .align 2
22    memory_start:
23        .word 0
24        .word 0
25        @ .word 1 @ I replaced this with 1 or 0 depending on the test I wanted to run.
26        .word 0
27        .word 0
28    memory_end:

```

Let's go through each point:

1. `.global _start` makes the entry symbol globally visible.
2. The addresses of `memory_start` and `memory_end` define the range that the program checks.
3. The registers `r0` and `r1` store the current address and the end address, respectively.
4. `ldr r2, [r0], #4` reads one 32-bit word into `r2` and advances the current address by four bytes.
5. `cmp r2, #0` and `bne explode` move execution to `explode` when the value is not zero.
6. `cmp r0, r1` and `bne check_memory` repeat the loop until the program reaches the end of the selected memory range.
7. The `all_zero` and `explode` labels use `b .` to identify the result and remain at the current instruction forever.
8. `.data` places the test values in the data section, while `.align 2` aligns them to a four-byte boundary.
9. The commented `.word 1` line lets me switch between the two cases by commenting or uncommenting it before rebuilding the program.

Now, it is about time to run it!

Running the all-zero test with GDB:

**Terminal 1: All-zero test**

```

→ ricardo homework1 arm-none-eabi-as -g memory_check.s -o memory_check.o
→ ricardo homework1 arm-none-eabi-ld -g memory_check.o -o memory_check.elf
→ ricardo homework1 qemu-arm -one-insn-per-tb -g 1234 memory_check.elf
(Waiting for connection...)

```

**Terminal 2: All-zero test**

```

→ ricardo homework1 arm-none-eabi-gdb
(gdb) file memory_check.elf
(gdb) target remote localhost:1234
(gdb) break all_zero
Breakpoint 1 at 0x801c: file memory_check.s, line 15.
(gdb) break explode
Breakpoint 2 at 0x8020: file memory_check.s, line 18.
(gdb) continue
Breakpoint 1, all_zero () at memory_check.s:15
15          b .
(gdb) info registers r2 pc
r2          0x0          0
pc          0x801c      0x801c <all_zero>

```

Next, I tested a memory range containing a nonzero value.

I uncommented `.word 1`, rebuilt the program, and repeated the same debugging procedure.

**Terminal 1: Nonzero test**

```

→ ricardo homework1 arm-none-eabi-as -g memory_check.s -o memory_check.o
→ ricardo homework1 arm-none-eabi-ld -g memory_check.o -o memory_check.elf
→ ricardo homework1 qemu-arm -one-insn-per-tb -g 1234 memory_check.elf
(Waiting for connection...)

```

**Terminal 2: Nonzero test**

```

→ ricardo homework1 arm-none-eabi-gdb
(gdb) file memory_check.elf
(gdb) target remote localhost:1234
(gdb) break all_zero
Breakpoint 1 at 0x801c: file memory_check.s, line 15.
(gdb) break explode
Breakpoint 2 at 0x8020: file memory_check.s, line 18.
(gdb) continue
Breakpoint 2, explode () at memory_check.s:18
18          b .
(gdb) info registers r2 pc
r2          0x1          1
pc          0x8020      0x8020 <explode>

```

**i Extra Credit** Please share how hard/easy this was for you. I want to adjust workloads/labs based on the difficulty of this first attempt.

This assignment took me  $\sim 4$  full hours of work. It looked short (and I thought it was going to take me  $\sim 1$  hour), but reading documentation, manual pages, and carefully understanding why everything works takes time. Also, I wanted to provide as many details as I could, but I am unsure if the level of detail I provided here is enough or if I overexplained things here. I probably spent more time than required because I overexplained some things.

I really liked the format; it is more like a conversation between the professor/TA and the student, so it is open to creativity. Also, the material is very well explained.

Fortunately, I use Arch Linux so I didn't have to troubleshoot a lot.