

STAT514 HW1 Solution

1 Q.1

1.1 a.

```
library(faraway)
data(prostate)
model1 <- lm(lpsa ~ lcavol, data = prostate)

> model1 <- lm(lpsa ~ lcavol, data = prostate)
> summary(model1)

Call:
lm(formula = lpsa ~ lcavol, data = prostate)

Residuals:
    Min      1Q  Median      3Q     Max 
-1.67625 -0.41648  0.09859  0.50709  1.89673 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 1.50730   0.12194 12.36   <2e-16 ***
lcavol      0.71932   0.06819 10.55   <2e-16 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.7875 on 95 degrees of freedom
Multiple R-squared:  0.5394,    Adjusted R-squared:  0.5346 
F-statistic: 111.3 on 1 and 95 DF,  p-value: < 2.2e-16
```

R-squared: 0.5394, Adjusted R-squared: 0.5346

1.2 b.

```
model2 <- lm(lpsa ~ lcavol + lweight + svi + lbph + age + lcp + pgg45 + gleason,
              data = prostate)
```

```

> summary(modeT2)

Call:
lm(formula = lpsa ~ lcavol + lweight + svi + lbph + age + lcp +
    pgg45 + gleason, data = prostate)

Residuals:
    Min      1Q  Median      3Q     Max 
-1.7331 -0.3713 -0.0170  0.4141  1.6381 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 0.669337  1.296387  0.516  0.60693    
lcavol       0.587022  0.087920  6.677 2.11e-09 ***  
lweight      0.454467  0.170012  2.673  0.00896 **   
svi          0.766157  0.244309  3.136  0.00233 **   
lbph         0.107054  0.058449  1.832  0.07040 .    
age          -0.019637 0.011173 -1.758  0.08229 .    
lcp          -0.105474 0.091013 -1.159  0.24964    
pgg45        0.004525  0.004421  1.024  0.30886    
gleason      0.045142  0.157465  0.287  0.77503    
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.7084 on 88 degrees of freedom
Multiple R-squared:  0.6548, Adjusted R-squared:  0.6234 
F-statistic: 20.86 on 8 and 88 DF,  p-value: < 2.2e-16

```

R-squared: 0.6548, Adjusted R-squared: 0.6234

1.3 c.

The R-squared of the linear model with multiple predictors is larger than the one with R-squared with only one predictor. This means that the multiple predictors model's predictors account for a large portion of the variation in the dependent variable, suggesting a strong fit for the data.

1.4 d.

- Manual Computation in R

```

y <- prostate$lpsa

X <- model.matrix(~ lcavol + lweight + svi + lbph + age + lcp + pgg45 + gleason,
                  data = prostate)

beta_hat_manual <- solve(t(X) %*% X) %*% t(X) %*% y

```

```

> beta_hat_manual
[ ,1]
(Intercept) 0.669336698
lcavol      0.587021826
lweight      0.454467424
svi          0.766157326
lbph         0.107054031
age          -0.019637176
lcp          -0.105474263
pgg45        0.004525231
gleason      0.045141598

```

- Compare with lm()

```

model2 <- lm(lpsa ~ lcavol + lweight + svi + lbph + age + lcp + pgg45 + gleason,
              data = prostate)

coef(model2)

> coef(model2)
(Intercept) lcavol      lweight      svi          lbph
0.669336698 0.587021826 0.454467424 0.766157326 0.107054031
           age          lcp      pgg45      gleason
-0.019637176 -0.105474263 0.004525231 0.045141598

```

The manually estimated parameters using the matrix formulation of least squares are basically identical to those obtained from the lm() function.

1.5 e.

```

y <- prostate$lpsa

X <- model.matrix(~ lcavol + lweight + svi + lbph + age + lcp + pgg45 + gleason,
                  data = prostate)

n <- nrow(X); p <- ncol(X)

XtX_inv <- solve(t(X) %*% X)
beta_hat <- XtX_inv %*% (t(X) %*% y)

yhat <- X %*% beta_hat
res  <- y - yhat

RSS <- as.numeric(crossprod(res))
sigma2_hat <- RSS / (n - p)

```

```

var_beta <- sigma2_hat * XtX_inv
se_beta  <- sqrt(diag(var_beta))
t_vals   <- as.vector(beta_hat) / se_beta
df       <- n - p
p_vals   <- 2 * pt(abs(t_vals), df = df, lower.tail = FALSE)

alpha <- 0.05
out <- data.frame(
  term      = colnames(X),
  estimate  = as.vector(beta_hat),
  SE        = se_beta,
  t         = t_vals,
  df        = df,
  p_value   = p_vals,
  decision  = ifelse(p_vals < alpha, "Reject H0", "Fail to reject H0"),
  row.names = NULL
)

manual_tests <- subset(out, term != "(Intercept)")
print(manual_tests, digits = 4)

> print(manual_tests, digits = 4)
      term estimate    SE     t df p_value      decision
2  lcavol  0.587022 0.087920 6.6767 88 2.111e-09 Reject H0
3  lweight  0.454467 0.170012 2.6731 88 8.955e-03 Reject H0
4     svi  0.766157 0.244309 3.1360 88 2.329e-03 Reject H0
5    lbph  0.107054 0.058449 1.8316 88 7.040e-02 Fail to reject H0
6     age -0.019637 0.011173 -1.7576 88 8.229e-02 Fail to reject H0
7      lcp -0.105474 0.091013 -1.1589 88 2.496e-01 Fail to reject H0
8   pgg45  0.004525 0.004421  1.0235 88 3.089e-01 Fail to reject H0
9 gleason  0.045142 0.157465  0.2867 88 7.750e-01 Fail to reject H0

```

1.6 f.

```

model2 <- lm(lpsa ~ lcavol + lweight + svi + lbph + age + lcp + pgg45 + gleason,
             data = prostate)

print(coef(summary(model2)))

```

```

> print(coef(summary(model2))) # estimates, SE, t, p from lm()
   Estimate Std. Error    t value Pr(>|t|)
(Intercept) 0.669336698 1.296387471  0.5163091 6.069335e-01
lcavol      0.587021826 0.087920303  6.6767493 2.110698e-09
lweight     0.454467424 0.170012435  2.6731423 8.955363e-03
svi        0.766157326 0.244309148  3.1360157 2.328749e-03
lbph       0.107054031 0.058449214  1.8315735 7.039846e-02
age        -0.019637176 0.011172725 -1.7575995 8.229321e-02
lcp        -0.105474263 0.091013487 -1.1588861 2.496377e-01
pgg45      0.004525231 0.004421179  1.0235350 3.088604e-01
gleason    0.045141598 0.157464523  0.2866779 7.750328e-01

```

Part E and F shows the same result.

1.7 g.

```

y <- prostate$lpsa
X <- model.matrix(~ lcavol + lweight + svi + lbph + age + lcp + pgg45 + gleason,
                  data = prostate)

n <- nrow(X); p <- ncol(X)

XtX_inv <- solve(t(X) %*% X)
beta_hat <- XtX_inv %*% (t(X) %*% y)

res <- y - X %*% beta_hat
RSS <- as.numeric(crossprod(res))
sigma2_hat <- RSS / (n - p)

var_beta <- sigma2_hat * XtX_inv
se_beta <- sqrt(diag(var_beta))

t_crit <- qt(0.975, df = n - p)

lower <- as.vector(beta_hat) - t_crit * se_beta
upper <- as.vector(beta_hat) + t_crit * se_beta

ci_table <- data.frame(
  term = colnames(X),
  estimate = as.vector(beta_hat),
  lower95 = lower,
  upper95 = upper
)

```

```

> print(ci_table)
      term    estimate    lower95   upper95
(Intercept) (Intercept) 0.669336698 -1.906960983 3.245634379
lcavol       lcavol    0.587021826  0.412298699 0.761744954
lweight      lweight   0.454467424  0.116603435 0.792331414
svi          svi     0.766157326  0.280644232 1.251670420
lbph         lbph    0.107054031 -0.009101499 0.223209561
age          age     -0.019637176 -0.041840618 0.002566267
lcp          lcp     -0.105474263 -0.286344443 0.075395916
pgg45        pgg45   0.004525231 -0.004260932 0.013311395
gleason     gleason  0.045141598 -0.267786053 0.358069248

```

2 Q.2

2.1 a.

Minimize the residual sum of squares

$$S(\beta) = \|y - X\beta\|_2^2 = (y - X\beta)^\top (y - X\beta)$$

Differentiate and set the gradient to zero:

$$\nabla_\beta S(\beta) = -2X^\top (y - X\beta) = 0 \Rightarrow X^\top X \hat{\beta} = X^\top y$$

Since X has full column rank, $X^\top X$ is invertible, so the unique minimizer is

$$\hat{\beta} = (X^\top X)^{-1} X^\top y$$

(uniqueness follows because S is strictly convex: $X^\top X \succ 0$)

2.2 b.

Using $y = X\beta + \varepsilon$ and $E[\varepsilon] = 0$:

$$E[\hat{\beta}] = (X^\top X)^{-1} X^\top E[y] = (X^\top X)^{-1} X^\top (X\beta) = \beta$$

So $E[\hat{\beta}] = \beta$

2.3 c.

$$\begin{aligned}
\text{Var}(\hat{\beta}) &= \text{Var}((X^\top X)^{-1} X^\top y) \\
&= (X^\top X)^{-1} X^\top \text{Var}(y) X(X^\top X)^{-1} \\
&= (X^\top X)^{-1} X^\top (\sigma^2 I_n) X(X^\top X)^{-1} \\
&= \sigma^2 (X^\top X)^{-1}
\end{aligned}$$

Thus $\text{Var}(\hat{\beta}) = \sigma^2 (X^\top X)^{-1}$

2.4 d.

Write $\hat{\beta} = Ay$ with $A = (X^\top X)^{-1}X^\top$. Because $y \sim \mathcal{N}(X\beta, \sigma^2 I_n)$ and linear transforms of multivariate normals are normal,

$$\hat{\beta} \sim \mathcal{N}(A(X\beta), A(\sigma^2 I_n)A^\top) = \mathcal{N}(\beta, \sigma^2(X^\top X)^{-1})$$

So

$$\hat{\beta} \sim \mathcal{N}(\beta, \sigma^2(X^\top X)^{-1})$$

3 Q.3

3.1 a.

Define the $(p+1) \times (p+1)$ matrix

$$C = \begin{bmatrix} 1 & c_{10} & c_{20} & \cdots & c_{p0} \\ 0 & c_{11} & c_{21} & \cdots & c_{p1} \\ 0 & c_{12} & c_{22} & \cdots & c_{p2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & c_{1p} & c_{2p} & \cdots & c_{pp} \end{bmatrix}$$

Its first column is $(1, 0, \dots, 0)^\top$ so the intercept is unchanged; for $j = 1, \dots, p$ the $j+1$ -st column of C is

$$\begin{bmatrix} c_{j0} \\ c_{j1} \\ \vdots \\ c_{jp} \end{bmatrix}$$

which encodes $z_j = c_{j0} \cdot \mathbf{1} + \sum_{k=1}^p c_{jk} x_k$. Then column-by-column, $Z = XC$ because

$$Z_{\cdot,1} = X_{\cdot,1}, \quad Z_{\cdot,j+1} = X C_{\cdot,j+1} = c_{j0} \mathbf{1} + \sum_{k=1}^p c_{jk} x_k = z_j$$

3.2 b.

OLS predictions are $\hat{y}_X = H_X y$ and $\hat{y}_Z = H_Z y$ with

$$H_X = X(X^\top X)^{-1}X^\top, \quad H_Z = Z(Z^\top Z)^{-1}Z^\top$$

Using $Z = XC$ and invertible C ,

$$\begin{aligned}
H_Z &= XC(C^\top X^\top XC)^{-1}C^\top X^\top \\
&= XC(C^{-1}(X^\top X)^{-1}C^{-\top})C^\top X^\top \\
&= X(X^\top X)^{-1}X^\top \\
&= H_X
\end{aligned}$$

Hence $\hat{y}_Z = \hat{y}_X$. Equivalently, since C is invertible, $\mathcal{C}(Z) = \mathcal{C}(X)$; both models project y onto the same column space, so fitted values (and thus residuals, RSS, R^2) are identical.